

Building scalable and secure L2 and L3 overlays with Host Identity Protocol

Dmitriy Kuptsov

2025

Contents

Contents	3
1. Introduction	5
1.1 Questions	6
2. Background	9
2.1 Cryptography basics	10
2.1.1 Symmetric cryptography	10
2.1.2 Asymmetric cryptography	11
2.1.3 Cryptographic hash functions	12
2.1.4 Key exchange protocols	13
2.1.5 Post-quantum Lattice-based cryptography	13
2.2 Security protocols	15
2.2.1 Host Identity Protocol (HIP)	15
2.2.2 Transport Layer Security (TLS)	16
2.2.3 Secure Shell Protocol (SSH)	17
2.3 L2, L3 and L4 tunneling	17
2.3.1 Virtual Private LAN Services (VPLS) solutions	17
2.3.2 Virtual Private Network (L3-VPN) security solutions	18
3. Results	21
3.1 Hardware-enabled symmetric cryptography	21
3.2 Host Identity Protocol based VPLS	22
3.3 Scalable multipoint to multipoint VPN using HIP protocol .	28
3.4 Comparison of various solutions	29
4. Conclusions	33
Bibliography	35

1. Introduction

Back at the end of *1960's* when the Internet was a rather small network, which was interconnecting major universities, governmental and military organizations, very little attention was devoted to security. Nowadays, when the Internet has become extremely sophisticated in structure, connecting billions of devices ranging from small IoT-type devices to humongous data centers, security has gained number one priority. In present days, a typical Intranet of an organization can include a number of geographically separated branch-office networks (for example, consider a factory that has many SCADA devices and a mission control center that is miles and miles away). Since these networks are geographically separated, connecting them becomes a necessity, and so is the security of these networks. This is when the layer-3 virtual private networks (*L3-VPN*) and layer-2 virtual private LAN service (*L2-VPLS*) solutions become handy. Scalability, and resilience to various attacks, from man-in-the-middle to integrity violation attacks, to rather fundamental attacks on asymmetric algorithms (such as RSA, DSA, and their elliptic curve counterparts, Diffie-Hellman and Elliptic Curve DH, for example) using, for example, *Shor's quantum computer algorithm* to factorize large numbers, and massive brute force attacks on hash algorithms should be considered thoroughly. With this in mind, in this work, we present different security solutions, which can be used to build secure L2 and L3 overlay networks. We present the limitations of each solution and identify how they can be avoided.

We start with background material on cryptography. Here we discuss various symmetric and asymmetric encryption algorithms, present the definition of hash functions, which are considered secure nowadays, and discuss several key agreement algorithms. To make the discussion complete we present the threat that quantum computers pose for such algo-

rithms as RSA and DH, and discuss how post-quantum algorithms such as those that are based on the lattice can be used as an alternative to classical algorithms for encryption and signature constructions. Although not considered as part of the present work, future work can include the performance comparison of standardized RSA and DSA algorithms with the performance of lattice-based algorithms incorporated into for example Host Identity Protocol or even Transport Layer Security protocol. We then move on to a discussion of TLS, SSL, IPsec, HIP, and SSH protocols and how those can be used to achieve integrity and confidentiality of data transmitted over insecure channels. Afterward, we discuss the results we have obtained over several years. Here, we discuss our practical experience with scalable Host Identity Protocol-based L3-VPN and VPLS network which was built using the same protocol. We devote a separate section on hardware-accelerated versions of AES and SHA-256 algorithms. We conclude the results section with an analysis of the limitations of each solution and present the results for the various micro-benchmarking settings.

1.1 Questions

In this work, we ask several questions. These are not research questions, but rather practical questions that we try to answer to ourselves in order to understand the usability of Python-based security solution. Since our work focuses on the application of Host Identity Protocol (HIP) in VPN and VPLS settings we ask the following questions:

First, what is the performance of the pure Python-based implementation of symmetric key encryption and decryption routines as well as hash methods and how do they compare to implementation, which uses special AES and SHA-256 CPU instructions. Here our focus is on the micro benchmarking of two implementations of AES and SHA-256 hashing algorithms, identification of the bottlenecks and further recommendations for our prototype implementation of Host Identity Protocol based VPLS and L3-VPN.

Second, what is the scalability of Host Identity Protocol based VPLS and how does it perform in emulated environments such as Mininet. Here we seek the answer to the question of whether the HIP-VPLS is usable in environments close to real-life setups.

Third, what is the performance of Python-based HIP-VPLS on real hard-

ware. By asking such questions we want to find the application niche of our security solution. In addition, we elaborate on the practical configuration of HIP-VPLS using a central controller.

The final question relates to *to the deployment of scalable L3-VPN based on Host Identity Protocol.* Here we focus on rather a different approach to building secure networks: we consider L3-VPN where nodes in different branch offices form separate broadcast and multicast domains, but still can communicate with each other (with the assistance of IPv4 or IPv6 routing protocols). Here, we want to answer how to tackle the scalability issues of VPN network by adding hierarchy into the architecture.

2. Background

Since we are going to discuss the security protocols in this work, we begin this section with a shallow dive into cryptography basics. Here, we discuss symmetric and asymmetric cryptography algorithms, to make the description a little bit complete we show how the RSA algorithm works, discuss Diffie-Hellman (DH) and its Elliptic Curve counterpart. We should mention that the current understanding inside the cryptographic community is such that Shor's algorithm and its quantum computer implementation theoretically can efficiently factorize big numbers and solve discrete logarithm problems without trouble. This algorithm, if powerful enough quantum computers will exist shortly, puts the RSA and DH algorithms - the major building blocks of modern security solutions - at risk of being cracked (once the modulus of the RSA algorithm factorized into prime components, the private key of the RSA the algorithm can be easily recovered). We will conclude this part of the background material with the discussion of **post-quantum** computer public key encryption solution based on lattice (more specifically we will discuss Learning With Errors (LWE) the problem, which is at the heart of modern public key cryptography). We believe that, eventually, this type of cryptography will be the replacement for traditional RSA and DH algorithms, which rely on the hardness of factorization of the big numbers and discrete logarithm problems. In the epilogue of this section, we will put a few words on how lattice public key cryptography can be used, for example, together with Host Identity Protocol.

In the second part of the background material, we will review the basics of the Host Identity Protocol, Transport Layer Security Protocol, and Secure Shell Protocol, since these protocols are essential for understanding the secure tunneling protocols that we discuss in this work.

We will finalize the discussion of the background material with a short

overview of various L2, L3 and L4 tunneling solutions, including L2 802.1Q QinQ tunneling, L3 Multi-Protocol Label Switching (MPLS), L4 tunneling using TLS and SSH protocol.

2.1 Cryptography basics

Cryptography comes in many flavors: symmetric key cryptography (3DES, AES, Twofish, RC4) which, in turn, can be categorized into block cipher and stream cipher and asymmetric key cryptography (such as RSA, DSA, ECDSA). There are also key exchange protocols such as Diffie-Hellmann and Elliptic Cryptography DH for negotiation of common keys over insecure channels. Different algorithms applicable in different settings depending on requirements. Typically, as we will discuss later, symmetric key cryptography is used to protect data-plane traffic in networks, whereas, asymmetric-key cryptography is more applicable to the common key negotiation, authentication and identification purposes [20].

2.1.1 Symmetric cryptography

We start with the symmetric key cryptography. Common key and rather trivial operations such as permutations and substitutions are at the heart of any symmetric key cryptography algorithm. Although this type of cryptography is efficient because of the usage of efficient operations, it comes with a limitation though. In symmetric key cryptography, both sender and receiver need to share the same key, which complicates such important aspects as key distribution and revocation and so alone this encryption solution a very hard to use in modern cryptosystems. Typically, asymmetric key cryptography such as RSA or DH is used to derive session keys – TLS, HIP, and many other protocols follow this design idea.

Symmetric key cryptography comes in two different flavors: block and stream. For example, block cipher (such as AES, 3DES, Twofish [20]) use blocks of data (typically, the size of the key is 128, 192, 256 bits [20], and typical block size is 64, or 128 bits), and encrypts or decrypts one block at a time. There are different modes of operation, though, for block ciphers, examples are counter mode and cipher block chaining. The latter uses a so-called initialization vector to add extra randomness into the encryption process, and encryption of proceeding blocks depends on the output of the previous block. Modes of operations are important for security reasons.

However, not all modes of operation are useful and secure. For example, Electronic Code Book (ECB), while achieving fast processing and parallelization, is considered insecure in many settings.

The other type of symmetric key algorithm is stream cipher. Here the encryption and decryption are performed on separate bits, one bit at a time. CR4 is an example of a stream cipher. Stream ciphers are extremely important in real-time processing, for example, Wi-Fi uses stream ciphers to encrypt the data plane traffic.

2.1.2 Asymmetric cryptography

Asymmetric key cryptography, in its simplest form, is brilliant in the age of computing. Guessing from the name that this type of cryptography uses different keys for encryption and decryption does not require deep thought. This property makes this group of algorithms suitable for various key distribution, revocation, and signature ideas.

There is a magnitude of different asymmetric key security algorithms. RSA, DSA, and its Elliptic curve variant ECDSA are the pillars of modern security solutions. However, the flexibility of these schemes comes at an extra price of CPU cycles. All this makes these solutions inapplicable for securing data plane traffic, but only rather to secure control plane. In what follows, just to underpin the beauty of the math behind asymmetric key cryptography, we provide a description of the RSA algorithm.

In the RSA cryptosystem, the sender generates a pair of keys as follows: First, the sender chooses large enough two prime numbers p and q . Next, the sender computes $n = pq$ and evaluates Euler's phi function: $\phi(n) = (p - 1)(q - 1)$. This is the same as the number of numbers co-prime to n . The sender then selects at random encryption exponent e such that $1 < e < \phi(n)$ and also e should be co-prime to $\phi(n)$. Finally, the sender or the dealer computes the decryption exponent d , such that $ed \equiv 1 \pmod{\phi(n)}$ using modular multiplicative inverse (for that purpose extended Euclidean algorithm can be used).

The public key is then (n, e) , and the private key is (n, d) . To encrypt the message m the sender computes $c = m^e \pmod{n}$. The decryption is similar $m = c^d \pmod{n}$. The beauty is in Fermat's little theorem, which states that $m^{\phi(n)} \pmod{n} \equiv 1 \pmod{n}$. Now, $ed \equiv 1 \pmod{\phi(n)}$, which means that $ed = k\phi(n) + 1$, and so $m^{(ed)} \pmod{n} \equiv m^{(k\phi(n)+1)} \pmod{n} \equiv 1^k m \pmod{n} \equiv m \pmod{n}$.

In practice, RSA requires random padding to protect against such attacks as chosen ciphertext attacks and making two identical plaintexts

produce various ciphertexts. Padding also ensures that the message size is multiple of the encryption block-size. In practice, Optimal Asymmetric Encryption Padding (OAEP) scheme is used.

It is good to know that if the message is hashed and encrypted with the private key, the result is a form of digital signature since the sender cannot later deny that it was involved in the encryption process. A Digital Signature Algorithm (DSA) is another example of an asymmetric signature scheme and was specifically designed for that purpose. In turn, Elliptic Curves improve the performance of regular DSA algorithms.

Frankly speaking, one-way functions can be also used to construct signature schemes. For example, one can use one-time hash-based signatures to produce secure digital signatures. Nevertheless, the application of these types of signature algorithms is rather impractical and finds little application in real-life settings.

2.1.3 Cryptographic hash functions

Mathematically, speaking hash function is a special one-way function: For a given pre-image of an arbitrary size it produces an image or hash value of a fixed size, which is universally unique. Ideally, secure hash functions should guarantee that the result it produces is irreversible. That is, it should be extremely hard to find a pre-image or original message, given the hash or the fingerprint. Secure hash functions should be also collision-resistant. In other words, it should be extremely hard, if not impossible at all, to find two different messages m and m' that will hash to the same value, i.e., $\text{hash}(m) = \text{hash}(m')$.

Secure hash functions are important in modern cryptography. For example, they can serve as authentication tokens for messages transmitted over the wire (useful, for example, in detecting message manipulation during transmission), they also allow compressing the message before signing it with the digital signature algorithm, and, finally, they can be used to find the differences between the messages efficiently (useful in large file transfer operations). The application area is of course broader than just these few examples.

Hash functions come in different flavors, but good ones should be computationally efficient and resistant to collisions. Today, hash functions such as MD2, MD4 and MD5 considered broken, as there are works that showed successful attacks. Briefly speaking, researchers found collisions for these hash functions. Therefore, it is not recommended to use these

hash functions in security applications. A more modern family of SHA hash functions also exists. For example, engineers recommend to use SHA-256, SHA-512 and recent SHA-3 in modern applications, as no successful attacks were registered for these types of hash functions.

Hash functions pave the road for such a notion as authentication tokens when combined with a secret key in a special way. Examples are Hash-based MAC (HMAC) [20], Parallelizable MAC (PMAC) [14], Cipher-based MAC (CMAC) which is based on AES cipher. For instance, by sending an HMAC together with the original message one can make sure that the message will not be modified during the transmission. If, however, the message will be altered on the route to a recipient, this fact will be detected immediately during the verification process.

Hash functions are also useful in signatures. For example, one-time signatures use hash functions to construct a digital signature of a message. They are, however, impractical as they require a considerable amount of storage and can be used only one time as the name implies. An interested reader can find more information about hash functions here [20].

2.1.4 Key exchange protocols

Finally, key exchange algorithms are also important in modern systems as they allow the negotiation of common keys over insecure channels. Of course, RSA can be used to deliver a session key by encrypting it with the recipient's public key, but specially crafted key negotiation algorithms exist in practice. Two bright examples are Diffie-Hellman (DH) and Elliptic Curve DH. Both DH and ECDH need to be authenticated in order to guarantee security.

2.1.5 Post-quantum Lattice-based cryptography

Shor's algorithm [18], implemented on a quantum computer, makes certain computational problems (such as factorization of large numbers and discrete logarithm problems) feasible in polynomial time. This shutters the security of the Internet, and so rigorous research was initiated to fill the gap. In what follows we discuss certain hard mathematical problems on lattices and show the workings of the Learning With Errors (LWE) public key encryption scheme [17]. In fact majority of NIST's candidates for post-quantum public key encryption algorithms are based on LWE.

A lattice is a mathematical structure that consists of integers in n di-

mensions arranged in a structured lattice-like way. Mathematically, the lattice is defined as follows:

$$\Lambda(\mathbf{B}) = \{\mathbf{B}\mathbf{x}, \mathbf{x} \in \mathbb{Z}^n\}$$

where \mathbf{B} is a matrix of basis vectors that generates the lattice. We should note that there exist a large number of basis vectors, some are *good* some are *bad*.

A **closest vector problem (CVP)** on lattices, which is considered NP-hard, and believed unsolvable even on quantum computers, can be defined as follows. Given a point $t \in \mathbb{R}^n$ and a lattice $\Lambda(\mathbf{B})$, the task is to find a closes point \mathbf{Bx} on lattice:

$$\min_{\forall \mathbf{x} \in \mathbb{Z}^n} \|\mathbf{B}\mathbf{x} - t\|$$

In practice, the above problem is extremely hard to solve which makes lattice-based cryptography attractive to cryptographers.

From linear algebra we know that solving equation $\mathbf{Ax} = \mathbf{b}$ is simple using Gaussian elimination. However, if a random noise is added to the equation

$$\mathbf{Ax} + \mathbf{e} = \mathbf{b}$$

the problem is considered as hard as CVP on the lattice. Solving the above problem directly relates to solving the CVP problem on lattice if the parameters are selected carefully.

So, given a matrix $\mathbf{A} \sim \mathbf{U}(\mathbb{Z}_q^{nxm})$, vector $\mathbf{s} \sim \mathbf{U}(\mathbb{Z}_q^n)$ and vector $\mathbf{e} \sim \mathbf{D}_{\mathbb{Z}^m, \sigma}$ sampled from discrete (clipped) Gaussian distribution with parameter σ . We require that, the probability $P[e < q/4]$ is high (*i.e.* 99.99%) to ensure correct decryption of the message and to achieve the required level of security. We can define matrix \mathbf{A} , secret key \mathbf{s} and noise vector \mathbf{e} as follows:

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} & \dots & a_{1n} \\ a_{21} & a_{22} & a_{23} & \dots & a_{2n} \\ \dots & \dots & \dots & \dots & \dots \\ a_{m1} & a_{m2} & a_{m3} & \dots & a_{mn} \end{bmatrix} \quad (2.1)$$

$$\mathbf{s} = \begin{bmatrix} s_1 \\ s_2 \\ \dots \\ s_n \end{bmatrix} \quad (2.2)$$

$$\mathbf{e} = \begin{bmatrix} e_1 \\ e_2 \\ \vdots \\ e_m \end{bmatrix} \quad (2.3)$$

Once the parameters are generated, we can compute $\mathbf{As} + \mathbf{e} = \mathbf{b}$. Then, the public key is (\mathbf{A}, \mathbf{b}) and the private key is \mathbf{s} . Deriving \mathbf{s} from \mathbf{b} is a hard task at hand.

To encrypt the message $\mu \in \{0, 1\}$, we choose $\mathbf{r} \sim U(\{0, 1\}^m)$. Then we compute $\mathbf{u} = \mathbf{r}\mathbf{A}$ and $v = \mathbf{rb} + \lfloor q/2 \rfloor \mu$. The ciphertext is (\mathbf{u}, v) . To decrypt the message we can compute $v - \mathbf{us}$: if the result is less than $q/4$ output 0, otherwise, if the result is larger than $q/4$ output 1. For decryption to work correctly, we require that the parameter $\sigma = q/(16m)$.

The major disadvantage of lattice-based cryptography is the size of the keys and actual ciphertext. For example, the security of LWE depends on two parameters n and q . By choosing $n = 512$ and $q = 2^{16}$, the size of ciphertext for a message of $k = 256$ bits long (for example, this is the size of the key for AES-256 symmetric algorithm), will be $O(k \cdot n \cdot \log q) \approx 256 \cdot 16 \cdot 512$ bits or roughly whooping 256 KB. All in all the security does not come for free. Of course, there are way much practical implementations of LWE-based encryption algorithms, for example, the reader can take a look at Kyber [16] which has a practical implementation in the TLS library.

2.2 Security protocols

Equipped with a basic understanding of cryptography we will now dive into a discussion of some of the well-known security protocols, including IPSec, HIP, TLS, and SSH. All these protocols make a solid basis for secure internetworking.

2.2.1 Host Identity Protocol (HIP)

Internet was designed initially so that the Internet Protocol (IP) address has a dual role: it is the locator, so that the routers can find the recipient of a message, and it is an identifier so that the upper layer protocols (such as TCP and UDP) can make bindings (for example, transport layer sockets use IP addresses and ports to make connections). This becomes a problem when a networked device roams from one network to another, and so the IP address changes, leading to failures in upper-layer connections. The

other problem is the establishment of an authenticated channel between the communicating parties. In practice, when making connections, the long-term identities of the parties are not verified. Of course, solutions such as SSL can readily solve the problem at hand. However, SSL is suitable only for TCP connections, and most of the time, practical use cases include only secure web surfing and the establishment of VPN tunnels. Host Identity Protocol, on the other hand, is more flexible: it allows peers to create authenticated secure channels on the network layer, so all upper-layer protocols can benefit from such channels. More on the protocol can be found in [15].

HIP relies on the 4-way handshake to establish an authenticated session. During the handshake, the peers authenticate each other using long-term public keys and derive session keys using Diffie-Hellman or Elliptic Curve (EC) Diffie-Hellman algorithms. To combat the denial-of-service attacks, HIP also introduces computational puzzles.

HIP uses a truncated hash of the public key as an identifier in the form of an IPv6 address and exposes this identifier to the upper layer protocols so that applications can make regular connections (for example, applications can open regular TCP or UDP socket connections). At the same time, HIP uses regular IP addresses (both IPv4 and IPv6 are supported) for routing purposes. Thus, when the attachment of a host changes (and so does the IP address used for routing purposes), the identifier, which is exposed to the applications, stays the same. HIP uses a particular signaling routine to notify the corresponding peer about the locator change. More information about HIP can be found in RFC 7401 [3].

2.2.2 Transport Layer Security (TLS)

Secure socket layer (SSL) [2] and Transport Layer Security (TLS) [5] are an application layer solutions to secure TCP connections. SSL was standardized in RFC 6101. TLS was standardized in RFC 5246. And was designed to prevent eavesdropping, man-in-the-middle attacks, tampering, and message forgery. In SSL communicating hosts can authenticate each other with the help of longer-term identities - public key certificates. SSL is great for building VPN tunnels and protecting upper-layer protocols such as HTTP.

2.2.3 Secure Shell Protocol (SSH)

Secure Shell protocol (SSH) is the application layer protocol that provides an encrypted channel for insecure networks. SSH was originally designed to provide secure remote command-line, login, and command execution. But in fact, any network service can be secured with SSH. Moreover, SSH provides a means for creating VPN tunnels between spatially separated networks: SSH is a great protocol for forwarding local traffic through remote servers.

2.3 L2, L3 and L4 tunneling

Virtual Private LAN Services (or VPLS), L3-VPNs, and L4 tunneling are pretty standard nowadays. Companies build security solutions to provide Layer-2 and Layer-3 services for branch offices: VPLS are typically built as overlays on top of Layer-3 (IP) and are Ethernet over IP type overlays, whereas L3-VPNs are IP-in-IP tunneling solutions.

In VPLS, when a frame arrives at VPLS provider equipment (PE), it is encapsulated into an IP packet and is sent out to all other VPLS network elements comprising emulated LAN. Security of such overlays is important for obvious reasons: customers do not want their corporate traffic to be sniffed and analyzed. In L3-VPN networks, on the other hand, the networks form different broadcast domains, and so when an IPv4 or IPv6 packet arrives at the VPN box, it is encapsulated in another IP packet and sent out using the backbone network. In this work, we built such secure overlays with Host Identity Protocol.

In this section, however, we will briefly review some of the widely used solutions for building L2, L3 and L4 overlays.

2.3.1 Virtual Private LAN Services (VPLS) solutions

In this section we will cover to standard ways to build VPLS networks (using, for example, 802.1q QinQ tunneling and MPLS).

QinQ tunneling

When the path from one network to the other, such as branch office to head office, traverses only layer-2 switches (*i.e.* no IP routing is involved), the VPLS can be organized with the help of 802.1Q protocol [12]. Broadly, speaking this is not a protocol as such, but rather VLAN tag-based switch-

ing. Thus, on the ingress point, an additional 802.1q service provider SP-VLAN tag is inserted in the L2 header of an Ethernet frame. Later, the forwarding decisions are made using this SP-VLAN tag. On the egress point, the SP-VLAN tag is removed and the original Ethernet frame is forwarded to the recipient based on the destination MAC address and, if exists, on the inner C-VLAN tag.

It should be noted that the configuration of forwarding is a manual step. Also, QinQ does not provide additional mechanisms to secure the customer's traffic, thus limiting the application domain of this solution.

MPLS tunneling

Multi-protocol label switching is a standard protocol for forwarding any traffic type. It is a layer 2.5 solution that sits between the data link layer and the network layer.

In MPLS the packets are forwarded not using MAC or IP addresses, but rather using labels, which are distributed by control protocol. Thus, when a frame arrives at the router the current label is popped, the new label is added and the frame is forwarded to the next hop router. The process continues until the frame reaches the destination network where it is routed based on the original identifiers (IP addresses or MAC addresses). Obviously, MPLS has label distribution protocol and label switching components. MPLS is an ideal solution to create overlays (*i.e.* L2 and L3 VPNs).

2.3.2 Virtual Private Network (L3-VPN) security solutions

The major drawback of QinQ and MPLS is that they do not offer encryption and authentication of traffic out-of-the box. Therefore, additional steps needs to be taken to protect end-to-end traffic. In this section we will review PPTP, SSL-based VPNs, L2TP and IPsec tunnels.

Multipoint to single point VPN

Multipoint to single-head VPN is a standard way of organizing a VPN network for an organization that has a single head office and multiple branch offices. In this setup, multiple branch offices are connected to a head end. We show such a setup in Figure 2.1.

There are several protocols available for such an arrangement. Examples are: (i) Point-to-Point Tunneling Protocol (PPTP) [19]; (ii) Generic Routing Encapsulation (GRE) [19]; (iii) SSL-based Secure Socket Tunnel-

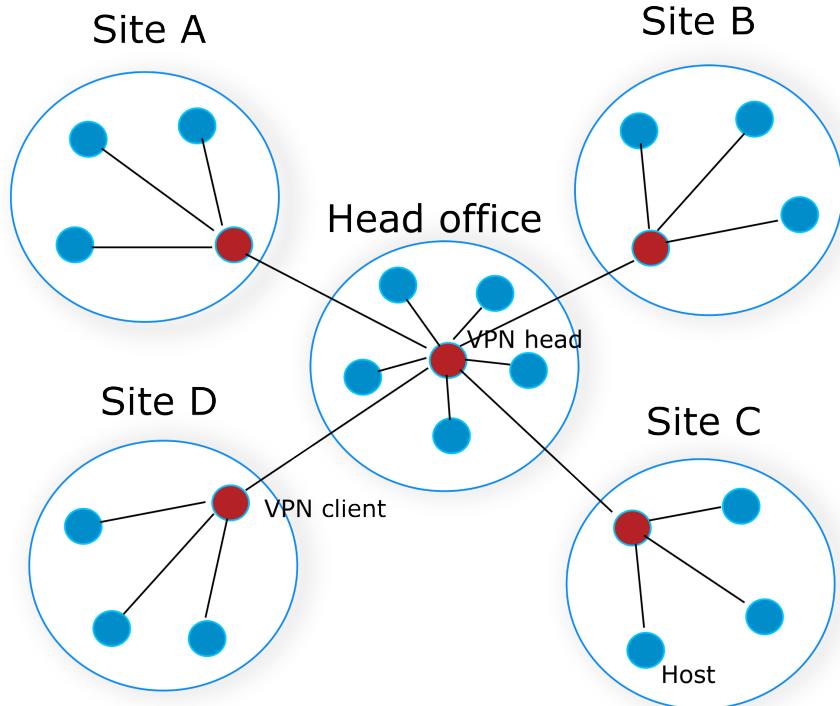


Figure 2.1. Typical arrangement of the VPN

ing Protocol (SSTP); (iv) Layer 2 Tunneling Protocol (L2TP) [19], which is an older protocol that can be combined with IPsec for encryption; (v) Internet Protocol Security (IPSec).

GRE on its own does not provide security and can be used together with IPsec to secure the traffic. PPTP, in turn, does not provide strong security out of the box. PPTP uses weak Microsoft Point-to-Point Encryption (MPPE), which is considered insecure. PPTP combines GRE and PPP protocols under a single umbrella. It is the PPP protocol [19] that provides such services as authentication (using MS-CHAPv2, PAP or strong EAP protocol) and link configuration (*i.e.* using Link Control Protocol (LCP)). Overall, it is not recommended to use PPTP in modern VPN setups. L2TP with IPsec AES-256 is a more secure alternative.

SSTP protocol is built on top of existing SSL. It allows tunneling user traffic over protected channel, and yet the traffic looks like normal HTTPS traffic to service providers. We have, ourselves, created a similar in spirit L3-VPN solution that is based on SSL [6]. The solution operates on a standard HTTPS port. However, our idea is to tunnel all traffic from VPN-agnostic hosts through the off-the-path black box that encrypts all traffic and sends encapsulated in TCP and SSL packets to the L3-VPN head server. The solution that we have created is a simple script that allows us to set up such an arrangement with no hassle. One drawback is that it

uses TCP for transport: sending over a reliable TCP channel and over well known HTTPS port is good for bypassing the traffic filters, but reduces the performance especially if the channel has a large latency and error rate.

IPSec [19] comes in two variations: Authentication Header (AH) and Encapsulating Security Payload (ESP). The first does not encrypt the data-plane traffic but rather adds HMAC to the packet. The second one, in addition to authentication, adds encryption of the payload. IPSec, when combined with the key exchange protocols, such as Internet Key Exchange (IKE) [4], can be used to create secure tunnels between the sites.

SSH tunneling

SSH, despite that it was invented for remote access to Linux-like boxes, can be used to tunnel local traffic to remote machine and remote traffic to local machine [13]. Thus, it can be used to create layer-4 tunnels. For example, the following command will tunnel all local traffic from port 4443 to remote web-server *youtube.com* on port 443:

```
ssh -L 192.168.1.1:4443:youtube.com:443 user@strangebit.io
```

In this example, when the client types `https://192.168.1.1:4443/` in the browser window, the traffic will be forwarded to the remote **youtube** server through the SSH server `strangebit.io`.

There is also a possibility to perform reverse tunneling, *i.e.* one can expose the local service to the world. For example, suppose you have a precious MySQL resource in your local network running on host 192.168.1.45 on port 3306, then you can expose the service to the world using the following command:

```
ssh -R 0.0.0.0:3306:192.168.1.45:3306 user@strangebit.io
```

This way various tunneling setups can be organized making SSH an attractive secure tunneling solution.

3. Results

In this chapter, we are going to present the results that we have obtained throughout the several years that we have spent building various systems. We start with the results for the cryptographic library which we have implemented to boost the performance of AES and HMAC algorithms on Intel CPUs. We then present the results for complete HIP-VPLS architecture and present the looking of the web interface which was used to configure the HIP switches. Finally, we present the design and implementation of the hierarchical L3-VPN in the Mininet emulator.

3.1 Hardware-enabled symmetric cryptography

Part of the work that we have done was related to porting parts of the code to pure C and special Intel CPU instructions. In this section, we will describe our achievements in this direction.

For the benchmarking, we have selected three implementations. The first one was pure Python based. For that purpose, we have used Py-Cryptodome library. The second implementation was a Python wrapper to the C library that used special Intel CPU instructions to boost the AES and SHA-based HMAC operations. The third implementation was pure C library which was using Intel NI instructions. The results for AES-256 and HMAC operations for varying block sizes are shown in Figure 3.1 and in Figure 3.2. The plots show the average running time in microseconds with the 95% confidence intervals.

What does this mean to HIP-VPLS performance? For a standard packet of size 1500 bytes we have compared the performance (combined HMAC and AES-256) and it turned out, on one hand, that the implementation of cryptography in pure C with special CPU instructions was 12.1 faster than pure Python implementation. On the other hand, Python implemen-

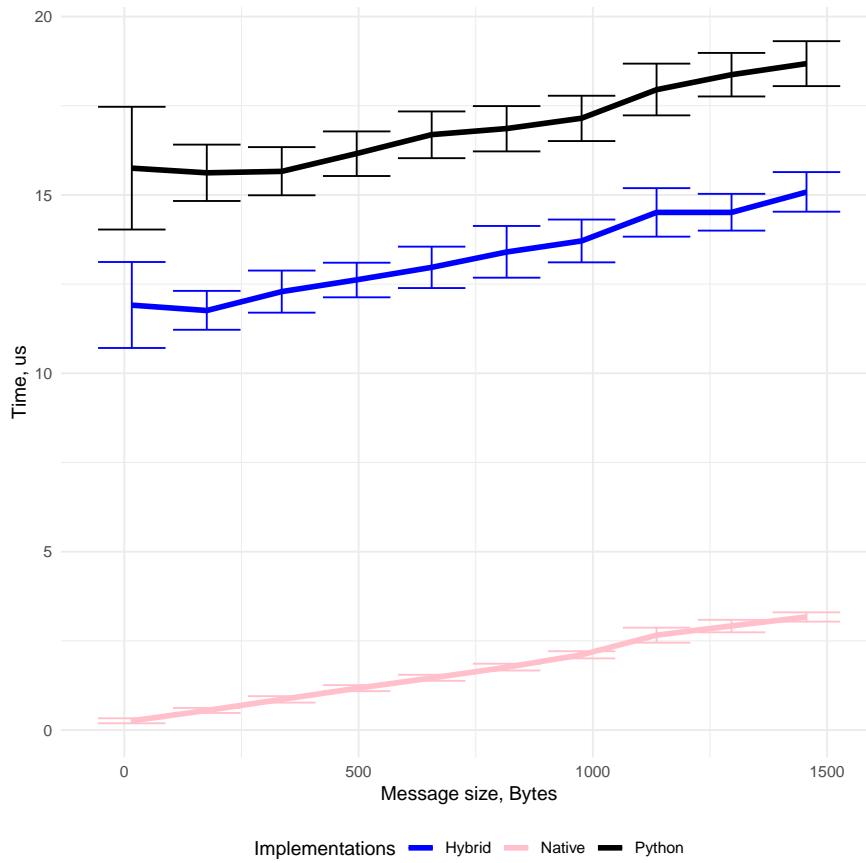


Figure 3.1. AES-256 encryption (microseconds)

tation with bindings to C library demonstrated performance which was 2.3 times faster. By making back of the envelop calculations we predict that Python implementation can achieve roughly 461 Mbit/s in upload and download directions cumulatively. However, in practice, given other operations with packets, we did not get this result in our experiments (more about the performance of HIP-VPLS on real hardware can be found in the proceeding chapter). For the plain C implementation with AES and SHA instructions, the performance will be better and constitute an astonishing 2.5 Gbit/s. If someone needs to run the code in production the entire code needs to be rewritten in plain C or Rust programming language for adequate performance.

3.2 Host Identity Protocol based VPLS

Virtual Private LAN Services (VPLS) provide means for building Layer 2 communication on top of existing IP networks. As we have mentioned already, VPLS can be built using various approaches. However, when building a production-grade VPLS solution one needs to have a clear picture of

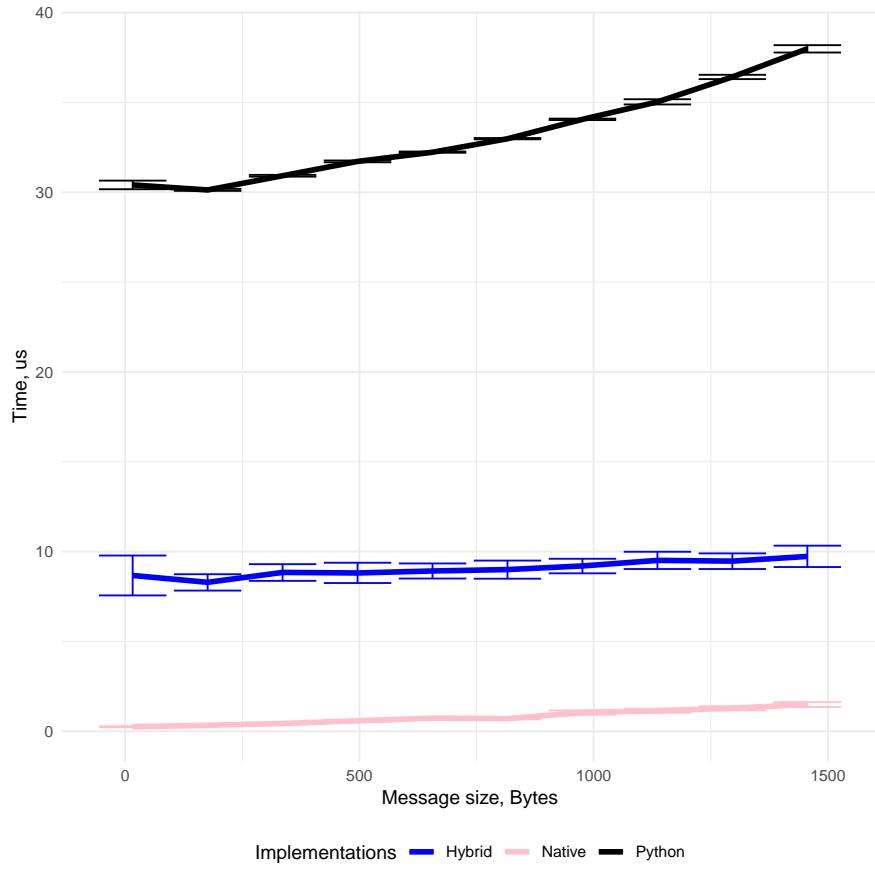


Figure 3.2. HMAC calculation (microseconds)

how such aspects as security and scalability will be solved.

In what follows, we will demonstrate how to build the VPLS using Host Identity Protocol (HIP). Our initial goal was not to build a production-grade implementation of HIP switches. Instead, at first, we were only interested in demonstrating proof of a concept solution in Mininet [1] – a framework for emulating L2 and L3 networks. It is worth mentioning that the code we have produced can be also deployed (under certain conditions; for example, our HIP implementation does not feature the NAT traversal mechanisms) on the real hardware in the Internet. We are going to demonstrate a working prototype in the later part of this work (here we assume that the public IPs are not from private range). All our prototypes use Python-based HIP [11] as the bases.

While building HIP switches (the switches that are deployed at the border of a network and are responsible for setting up security associations and pseudowires) we came across several challenges. First, to avoid loops the underlying network needs to support the IEEE 802.1D protocol (or its modification - this really depends on the version of the protocol supported by the switches). This problem was initially addressed in the rele-

vant IETF draft. For the sake of brevity, we note that if LAN implements 802.1D STP protocol there will be no loops in the HIP-VPLS instance. Second, there were certain issues with MTU and the inability of the Linux kernel to deliver IP packets when those are fragmented in user space and injected into the network stack using raw sockets. And finally, it took us some time to repackage the existing implementation of HIP protocol as a library, so that it would be agnostic about low-level networking (such as raw sockets, etc.). In the proceeding paragraphs, we will demonstrate the usage of HIP-based VPLS using loop-free L2 topology.

The logical network diagram of our Mininet prototype is shown in the Figure 3.4.

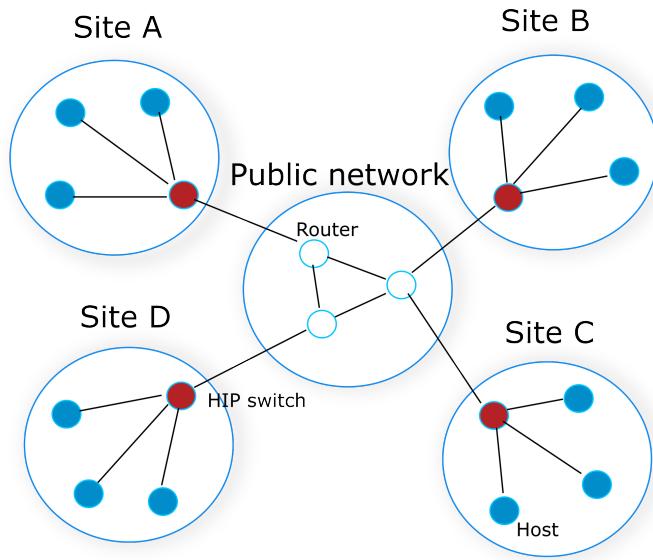


Figure 3.3. HIP-VPLS logical diagram (Mininet deployment)

Our HIP-VPLS implemnetation [7] in Mininet was using static configuration, meaning that HIP-VPLS mesh, resolver and firewall rules were configured prior to deployment of the overlay network and remained unchanged throughout the experiments. An interested reader can take a look at [7] for precise steps that are required to deploy the HIP-VPLS in the Mininet environment.

Overall, HIP-VPLS works as follows: (i) The daemon constantly listens for packets on private interface and public interface; (ii) if the frame, which arrives on the private interface, is broadcast or multicast daemon chooses all HIP-VPLS peers in mesh to send the packet; (iii) if the frame is unicast and HIP security association exists for the destination daemon sends the packet to the selected HIP switch; (iv) if no security association exists HIP switch triggers HIP base exchange to negotiate secret keys and to establish security association; (v) If the IPSec packet arrives on

the public interface, first HMAC is verified, and if it is valid the packet is decrypted; the original Ethernet frame is then reinjected into the private interface and regular destination MAC-based and VLAN-based forwarding is performed to deliver the frame to the recipient.

Our HIP switch also implements the MAC learning and aging functionality: whenever a frame arrives on the public interface the HIP-switch notes the source MAC address and adds it to the local database. Later, when a unicast frame arrives on the private interface, it looks up the destination MAC address and chooses the corresponding HIP association and pseudowire to send the frame encapsulated into an IPSec packet to the recipient.

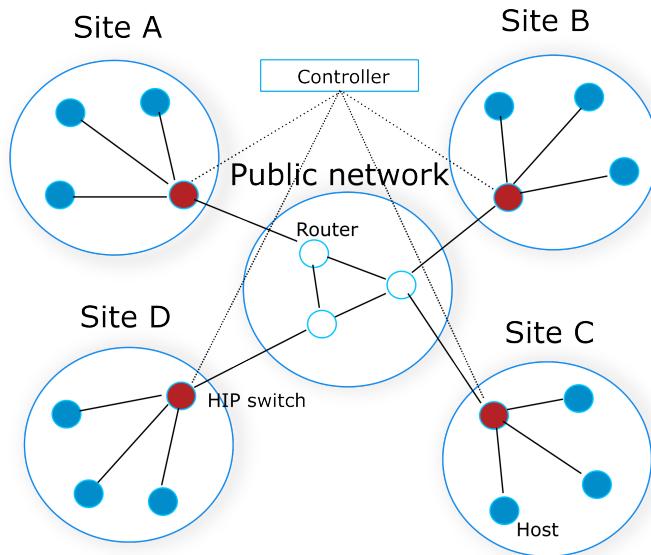


Figure 3.4. HIP-VPLS logical diagram (real hardware deployment)

To get a grasp on the performance of HIP-VPLS in the Mininet environment we have performed a series of bandwidth tests using *iperf* tool. To run the experiments we have used the UTM emulator (installed on MacBook M1) with Ubuntu 22.04 installed. All in all the results were the following: the 95% confidence interval for sample mean throughput was 58.9 ± 0.52 Mbit/s.

We now turn our attention to the real-life deployment of HIP-VPLS [8, 9]. The system architecture is similar to our Minent prototype (except that there was a lesser number of HIP switches) shown in Figure 3.4. Apart from the HIP-VPLS switches, we have also implemented a unique control-plane protocol on top of the SSL protocol for communication with the central controller on the Internet.

In our deployment, we have used the following setup. For HIP switches we have used the dual-network Intel N95 computing platform. We have

used 8 port SNR switch to connect 3 HIP switches, that way we have mimicked the IP overlay in the setup. HIP switches had two interfaces: one was facing LAN network, the other one was facing the WAN network.

The microcomputers for HIP switches had the following characteristics: they had 8GB of RAM memory, quad-core Intel N95 CPU (with support for AES and SHA2 NI instructions), 256 GB of solid state hard drive. To wire the routers we have used SNR switches (each switch had 8 1 Gbit/s ports and two Small Form Factor (SFP) slots). The testbed configuration is shown on Figure 3.6.

In the testbed, we had a multihomed server (with one IP facing the public network so that HIP switches will be able to connect to the controller in the Internet, and one IP in the private range; this server was playing the role of HIP controller), several legacy microcomputers, IP camera, and DHCP and DNS servers.

In our testbed the central controller was responsible reporting the liveness of HIP switches as well as provisioning the devices with the mesh configuration information, firewall rules and MAC-based ACL. For that purpose, we developed a simple secure protocol which was utilizing TLS. For example, consider the Figure 3.5 which shows the HIP switch registration and status information.

The screenshot shows a web-based interface titled "HIP-VPLS switch configurator". The top navigation bar includes links for "Home", "About", and "Logout". Below the navigation is a table listing five registered switches:

HIT	IP	Name	Last seen	Status
2001:0021:b097:0237:5bd6:6176:08cf:94ea	1.1.1.5	hip-switch-ohio.strangebit.io	Sun Jan 28 2024 19:52:15 GMT+0500 (Узбекистан, стандартное время)	Offline
2001:0021:f8e6:5867:4c14:2a78:c368:68d8	1.1.1.4	hip-switch-indiana.strangebit.io	Sun Jan 28 2024 20:59:32 GMT+0500 (Узбекистан, стандартное время)	Offline
2001:0021:e5b8:07c7:c47a:a469:5051:dd5c	1.1.1.3	hip-switch-idaho.strangebit.io	Sun Jan 28 2024 20:59:09 GMT+0500 (Узбекистан, стандартное время)	Offline
2001:0021:4b88:b52f:2563:c8e1:aa45:8e88	1.1.1.7	hip-switch-florida.strangebit.io	Mon Jan 29 2024 03:08:07 GMT+0500 (Узбекистан, стандартное время)	Online
2001:0021:6e40:3451:6726:acd5:5c4f:d043	1.1.1.8	hip-switch-massachusetts.strangebit.io	Mon Jan 29 2024 03:08:04 GMT+0500 (Узбекистан, стандартное время)	Online

Figure 3.5. HIP-VPLS central controller UI

According to the protocol, on the one hand, every HIP-VPLS the switch was reporting to the central controller (all requests were authenticated using the HMAC algorithm together with the shared symmetric master secret). In the implementation, switches were reporting their presence ev-

ery 5 seconds. On the other hand, every HIP-VPLS switch was obtaining the configuration from the central controller (such as mesh configuration, HIT resolver information, firewall rules, and MAC-based ACL).

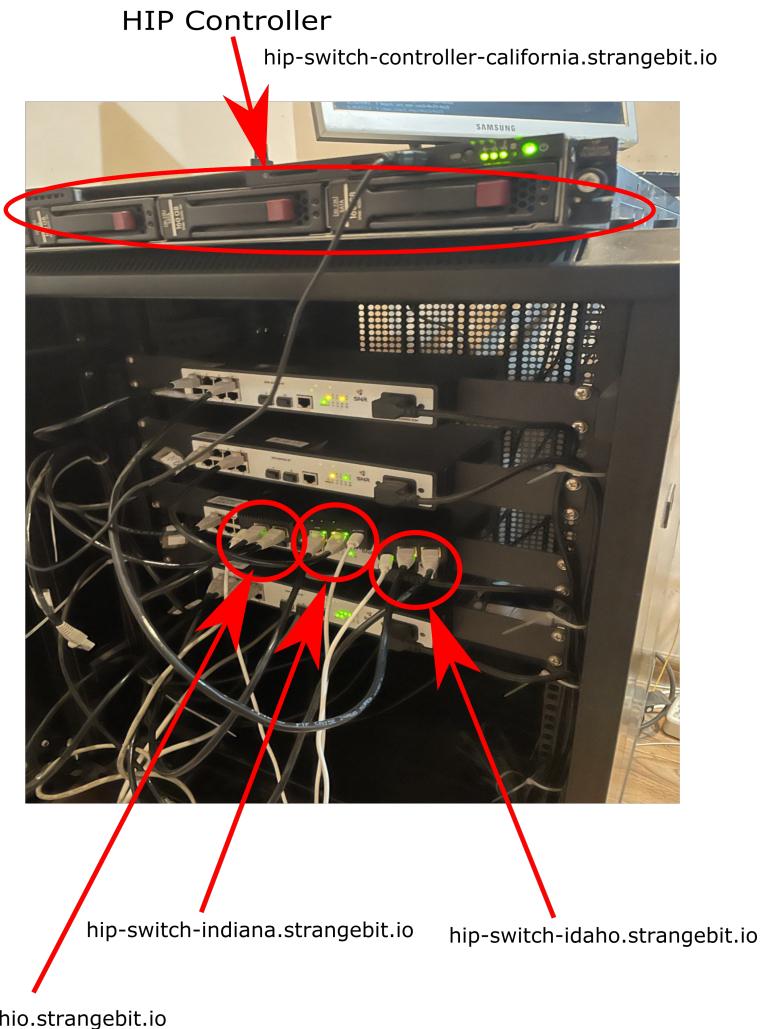


Figure 3.6. Testbed

To conclude we have performed a series of real-life experiments to measure the performance of the HIP-VPLS network. In Table 3.1 we show sample statistics for upload and download throughput. In addition, we have also measured latency. To perform the measurements we have used **speedtest** Python library. Thus, on one side, we have connected the MacBook to the HIP switch via a regular switch. On the other side, we have connected the other HIP switch to a network that had connectivity to the Internet. We then performed 100 rounds of measurements, collected throughput and latency data and processed the cleaned data using Python *statistics* library.

Statistics	Upload (Mbit/s)	Download (Mbit/s)	Latency (ms)
Sample mean	46.1	48.2	5.0
Sample std	7.1	2.3	0.19
Sample median	44.8	48.8	4.9
Sample min	14.3	40.0	4.6
Sample max	61.3	50.4	5.4

Table 3.1. Performance of HIP-VPLS on Intel N95 CPU

3.3 Scalable multipoint to multipoint VPN using HIP protocol

The major problem with the HIP-VPLS is the number of HIP switches and full-mesh connectivity between these switches. Imagine that there are not 10s, but 1000s sites, and that all sites need to be combined into a single network. First of all, there will be $O(n^2)$ pseudo-wires: for 1000 PEs there will be around 1M of routing table entries. Second, HIP-VPLS provides a single broadcast domain. And so there is going to be chaos in the network which will be overwhelmed with broadcast and multicast Ethernet frames. All these aspects make this type of arrangement of network unacceptable in the aforementioned scenarios. Instead, what if we let each site live in its own broadcast domain, *i.e.* have a separate network address, and combine through a series of overlay routers, which will be responsible for forwarding the packets between the networks (sites) based on inner IPv4 addresses.

To make the network scalable and reduce the number of pseudowires we let some nodes play the hub role, that is they will be the backbone of the overlay network. While some nodes will be the spoke nodes and will be connected directly to the sites. It is the hierarchy that makes the network scalable.

It is worth to look at the overall architecture which we have implemented in Mininet framework [10]. The logical diagram is shown in Figure 3.7. As we have already mentioned, the architecture of the distributed L3-VPN network is of hub-and-spoke type. Hub nodes comprise the backbone of the network, whereas, multiple spoke PE elements are attached to the hubs.

The security of the network is achieved by using Host Identity Protocol (on a hop-by-hop basis) to negotiate the authentication and encryption keys, whereas, the actual packet authentication and encryption is performed on hop-by-hop bases using HMAC-SHA256 and AES (with 256

bits key) algorithms. In our prototype implementation we have populated the routing tables manually, however, in practice this process should be automated using for example central controller.

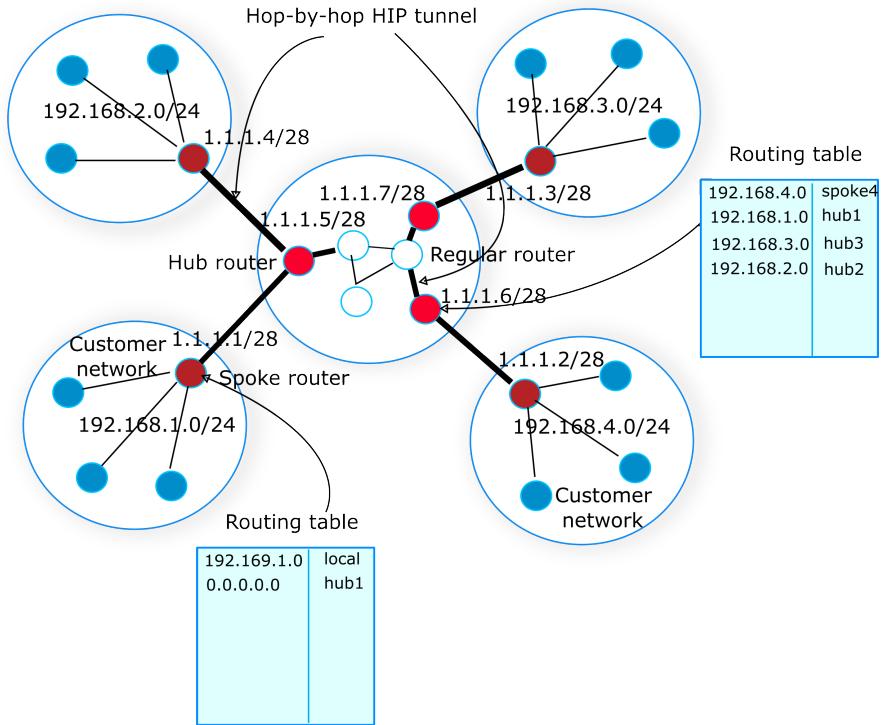


Figure 3.7. HIP-based L3-VPN in Mininet

To get the taste of the performance of this setup we have performed several rounds of experiments with the *iperf* utility and measured the throughput with encryption/authentication enabled. The results are the following: 19.7 ± 0.06 Mbit/s. This was expected, since the packet is decrypted and encrypted, as well as HMAC is recalculated at every hop on the path from source CE to destination CE. Perphas, hop-by-hop encryption and authentication can be done selectively with global secret key so that better performance can be achived.

3.4 Comparison of various solutions

In what follows, we compare now different approaches and identify their characteristics and limitations. In Table 3.2 we compare three different approaches for building overlays with Host Identity Protocol. The first one is scalable L2-VPLS with hub-and-spoke architecture. The second one, L3-VPN, also with hub-and-spoke design. And finally, we have HIP-VPLS at our disposal with full mesh connectivity of provider equipment (PE).

Characteristic ↓ Overlay type →	L2-VPLS	L3-VPN	HIP-VPLS
Size of forwarding/routing table	$O(n)$	$O(m)$	$O(n)$
Number of links in mesh	$O(k^2)$	$O(k^2)$	$O(l^2)$
Privacy (exposure of information)	MACs	IPs	No
Encryption and authentication	Hop-by-hop	Hop-by-hop	PE-to-PE
Tunneling mode	Ethernet-in-IP	IP-in-IP	Ethernet-in-IP
Loop free-topology	802.1D/Controller	Controller	Not required

Table 3.2. Comparison study of different multipoint VPLS/VPN designs

The first characteristic is the size of the forwarding table on the PE elements. For L2-VPLS and L3-VPN it is equal to $O(n)$, where n is the number of regular hosts in the network. This is obvious, as the MAC address table at least on the edge needs to know the mapping for each and every host in the network (consider when all hosts talk to all other hosts). For L3-VPN the size is considerably smaller since the routing table contains only IP prefixes of the networks and so equals to $O(m)$, where m is the number of sites, hence, the size of the network address prefixes. The reader should understand that $n \gg m$.

The second important characteristic is the number of links in a mesh network. For L2-VPLS and L3-VPN it is equal to $O(k^2)$, where k is the number of hub PEs. For HIP-VPLS this metric is equal to $O(l^2)$, such that l is the overall number of sites or PEs. Clearly, $l \gg k$, and hence the L3-VPN achieves better scalability.

What about privacy? Well in scalable L2-VPLS and L3-VPN the MAC and IPs are exposed to intermediate hubs (at the end these addresses are used for forwarding). And so if a hub gets compromised this information will be leaked to the adversary. In turn, in HIP-VPLS there are no intermediate nodes in the network since the pseudowires are created end-to-end, and so there is no risk that the customer will expose sensitive information to intermediate nodes. Also, in scalable L2-VPLS and scalable L3-VPN the encryption and authentication is done in a hop-by-hop manner; whereas, in HIP-VPLS the encryption is PE-to-PE (or site-to-site).

One last important point is the avoidance of loops in the network. For L2-VPLS loop-free topology is achieved with 802.1D protocol (PE should implement this functionality, because they perform forwarding tasks) or an SDN central controller. In L3-VPN the loops are avoided with the help of the IP TTL field. Also, in L3-VPN the routing tables are con-

structured centrally and no routing loops will exist in the topology. HIP-VPLS archives loop-free topology by assuming that customer networks run an instance of STP protocol. There is no need to implement 802.1D STP protocol for HIP switches since they do not forward Ethernet frames (received from public interface) to all, but private interface.

4. Conclusions

We started this work with the background material on cryptography. Here we covered established approaches (building blocks) of modern security protocols. In addition, we have introduced to the reader more recent developments, such as the LWE encryption scheme. We see that integration of LWE encryption and signature algorithm into HIP protocol can be future work. We then discussed how to build various secure tunnels, *e.g.* with SSL, IPSec, and SSH protocols. We covered briefly QinQ tunneling and MPLS protocol.

In the results section, we covered the results for various cryptographic libraries, including the library which uses Intel NI instructions designed to boost the AES and HMAC. We concluded that the Python library with C-bindings is not enough for the production setup, and suggested implementing the HIP-VPLS in Rust or C language. We then moved to the description of scalable L3-VPN and HIP-VPLS solutions. We concluded the work with a comparison of various characteristics of scalable L2-VPLS, L3-VPN and HIP-VPLS solution.

Bibliography

- [1] Mininet: An Instant Virtual Network on your Laptop. <https://mininet.org/>.
- [2] RFC 6101: The Secure Sockets Layer (SSL) Protocol Version 3.0. <https://datatracker.ietf.org/doc/html/rfc6101>.
- [3] RFC 7401: Host Identity Protocol Version 2 (HIPv2). <https://www.rfc-editor.org/rfc/rfc7401.html>.
- [4] RFC 4306: Internet Key Exchange (IKEv2) Protocol, <https://datatracker.ietf.org/doc/html/rfc4306>. Online, 2005.
- [5] RFC 5246: The Transport Layer Security (TLS) Protocol Version 1.2. <https://datatracker.ietf.org/doc/html/rfc5246>, 2008.
- [6] Bypassing Deep Packet Inspection: Tunneling Traffic Over TLS VPN <https://www.linuxjournal.com/content/bypassing-deep-packet-inspection-tunneling-traffic-over-tls-vpn>. Online, 2021.
- [7] Simulating Host Identity Protocol-Based Virtual Private LAN Service Using Mininet Framework. <https://www.linuxjournal.com/content/simulating-host-identity-protocol-based-virtual-private-lan-service-using-mininet-framework>, 2022.
- [8] HIP-VPLS. <https://github.com/strangebit-io/hip-vpls-hw-with-controller>, 2023.
- [9] HIP-VPLS controller. <https://github.com/strangebit-io/hip-vpls-controller>, 2023.
- [10] HIP-based L3-VPN. <https://github.com/dmitriykuptsov/hip-vpls-routing>, 2024.
- [11] Python-based Host Identity Protocol, <https://github.com/dmitriykuptsov/cutehip>. Online, 2025.
- [12] SNR S2980G-8T Switch Configuration Guide, <https://snr.systems/site/data-files/SNR%20Switches/Configuration%20Guide/SNR-S2980G-8T%20Configuration%20Guide%20v1.0.pdf>. Online, 2025.
- [13] SSH Tunneling, <https://www.ssh.com/academy/ssh/tunneling>. Online, 2025.

- [14] BLACK, J., AND ROGAWAY, P. A block-cipher mode of operation for parallelizable message authentication. In *Proceedings of the International Conference on the Theory and Applications of Cryptographic Techniques: Advances in Cryptology* (Berlin, Heidelberg, 2002), EUROCRYPT '02, Springer-Verlag, p. 384–397.
- [15] GURTOV, A. *Host Identity Protocol (HIP): Towards the Secure Mobile Internet*. 2008.
- [16] NATIONAL INSTITUTE OF STANDARDS AND TECHNOLOGY. FIPS 203: Module-Lattice-Based Key-Encapsulation Mechanism Standard, 2024. <https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.203.pdf>.
- [17] REGEV, O. The learning with errors problem (invited survey). In *Proceedings of the Annual IEEE Conference on Computational Complexity* (2010), pp. 191–204.
- [18] SHOR, P. Algorithms for quantum computation: discrete logarithms and factoring. In *Proceedings 35th Annual Symposium on Foundations of Computer Science* (1994).
- [19] STEVENS, W. R. *TCP/IP illustrated (vol. 1): The Protocols*. Addison-Wesley Longman Publishing Co., Inc., USA, 1993.
- [20] STINSON, D. *Cryptography: Theory and Practice, Second Edition*, 2nd ed. CRC/C&H, 2002.